

# Advanced Networking and Distributed Systems

## **Module 2: Scalable Servers and Network Performance**

GW CSCI 3907/6907

Timothy Wood and Lucas Chaufournier

# Outline

## Weeks 1-3: Network Programming and Protocols

- Writing simple network programs is easy!
- Providing reliable services over a network is hard!

## Weeks 4-5: Scalability and Performance

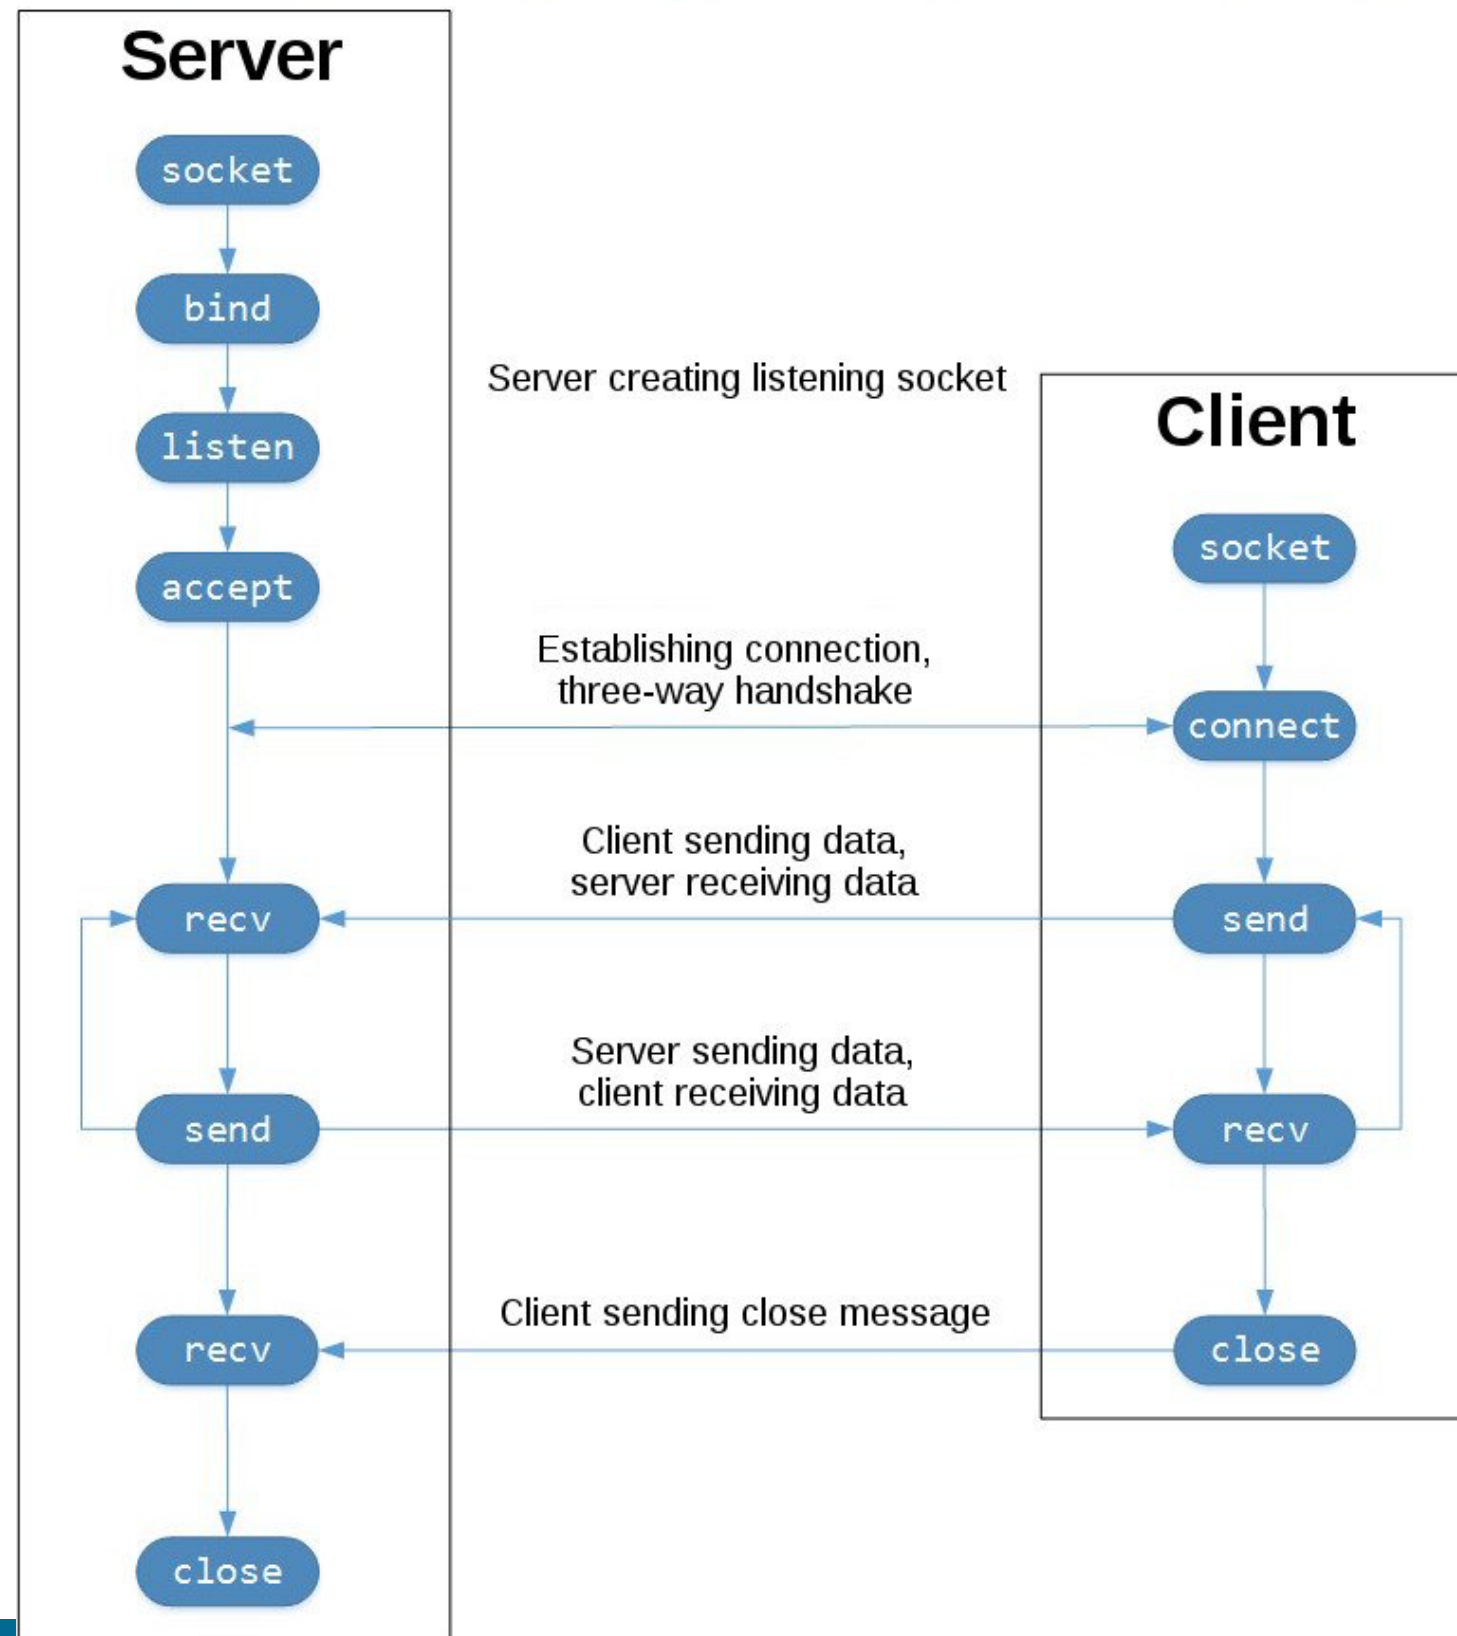
- How can we support many concurrent clients?
- What performance metrics matter for network services?

## Weeks 6-7: Network Middleboxes

- How to deploy software *between* clients and servers?
- How to get the speed of HW and flexibility of SW?

# Server Architecture

**How many clients  
can this server  
handle at once?**



# Simplest Architecture

Server is a single thread

Network calls are blocking (**recv**, **accept**)

-> server can only handle one client at a time

What happens to other clients who try to connect?

# Simplest Architecture


Server is a single thread

Network calls are blocking (**recv**, **accept**)

- Server can only handle one client at a time

What happens to other clients who try to connect?

- Incoming connections are buffered by the OS networking stack
- TCP Backlog parameter controls number of waiting connections
  - How do you think this works?



# Threading!



# Threading

Allows program to do multiple things at once

- Threads: execution context with its own stack and shared heap
- Processes: execution context with both stack and heap

How many threads or processes can we run?

# Threading

Allows program to do multiple things at once

- Threads: execution context with its own stack and shared heap
- Processes: execution context with both stack and heap

How many threads or processes can we run?

- Depends on available hardware and application type!

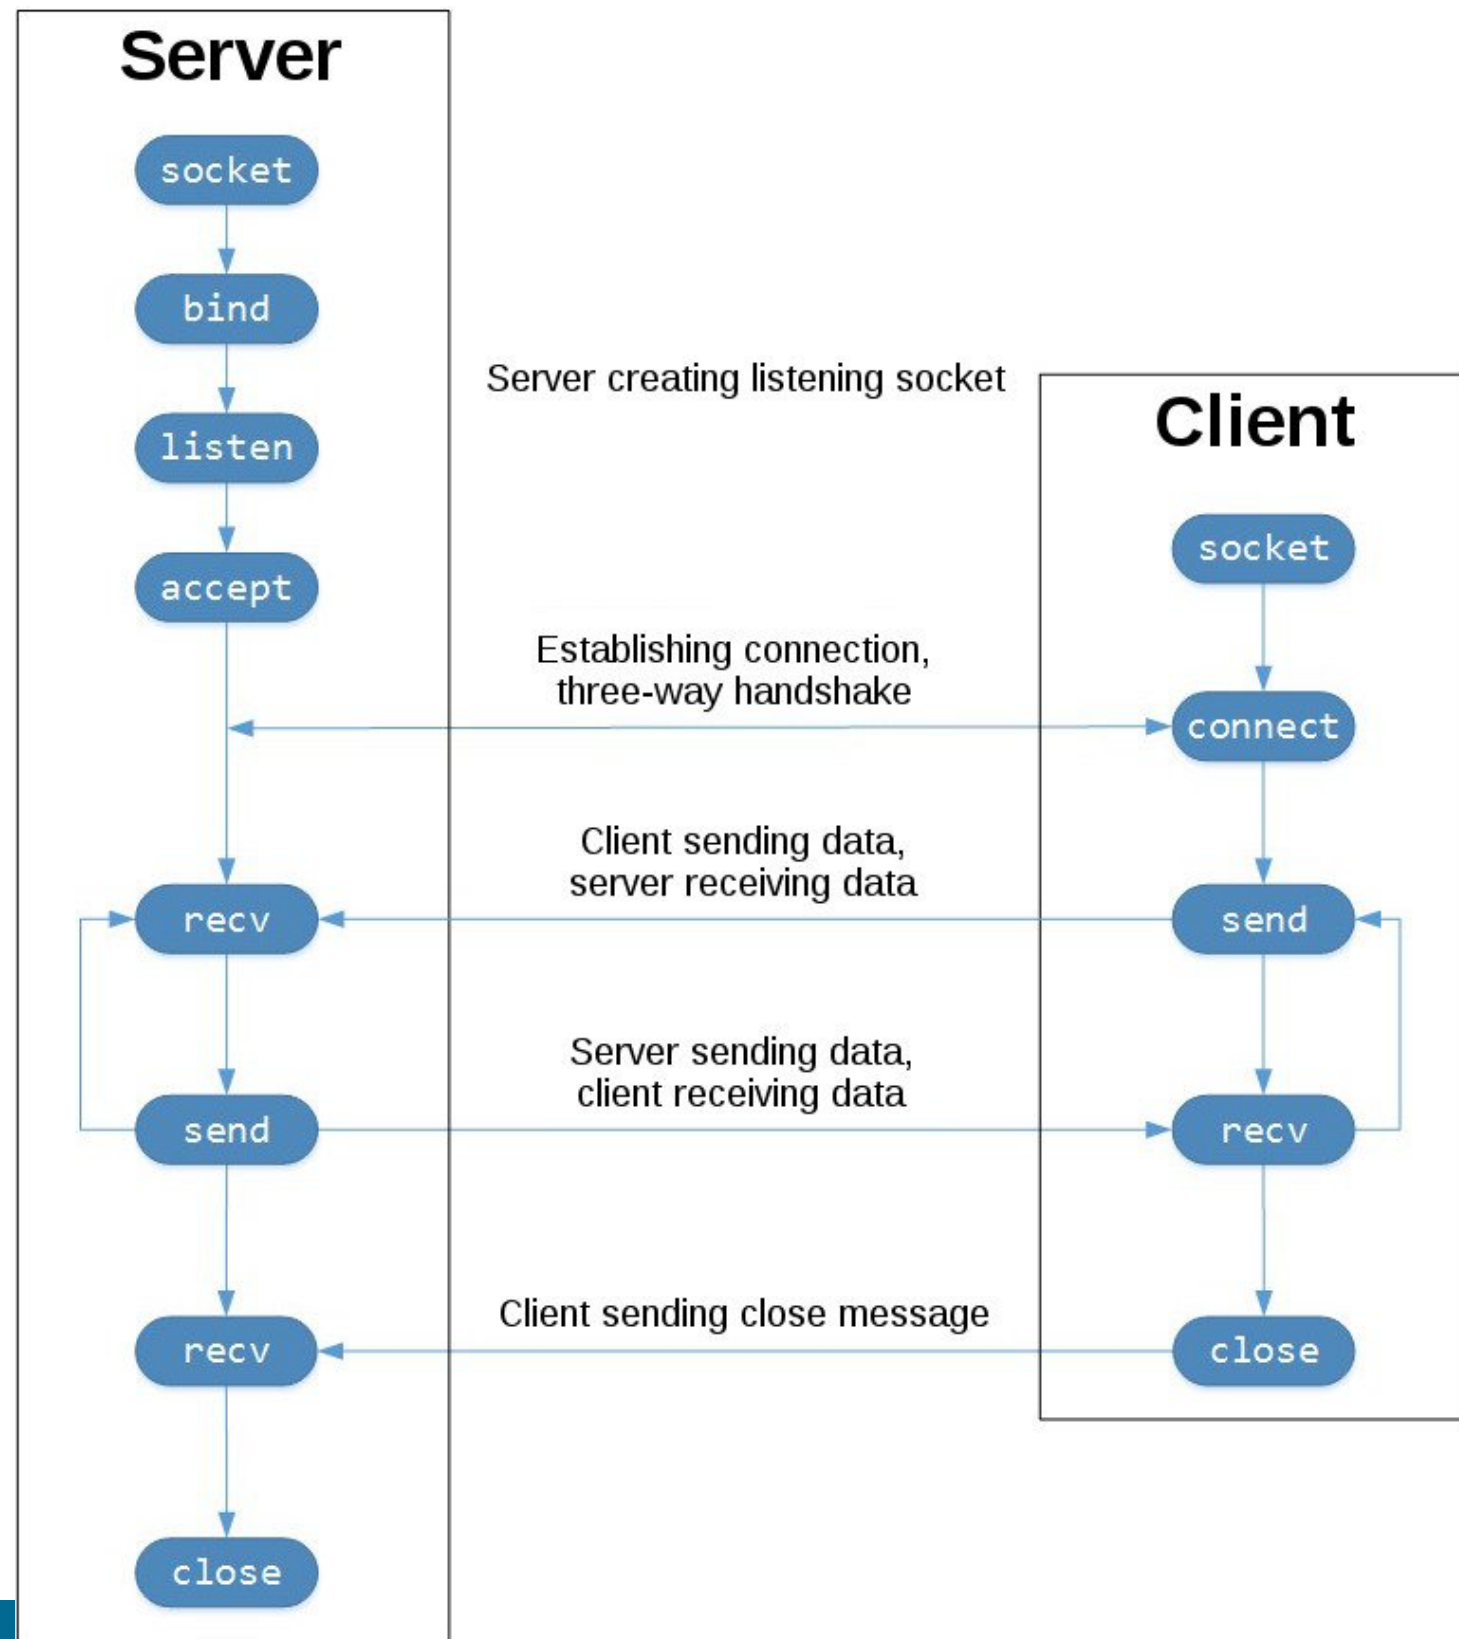
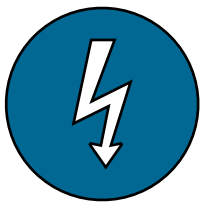
Concurrency is limited by...

- Number of CPU cores
- CPU vs IO intensiveness of application
- If CPU bound, then N cores can only run N threads at once
- If I/O bound, then may need  $\gg$  N threads to keep N cores busy



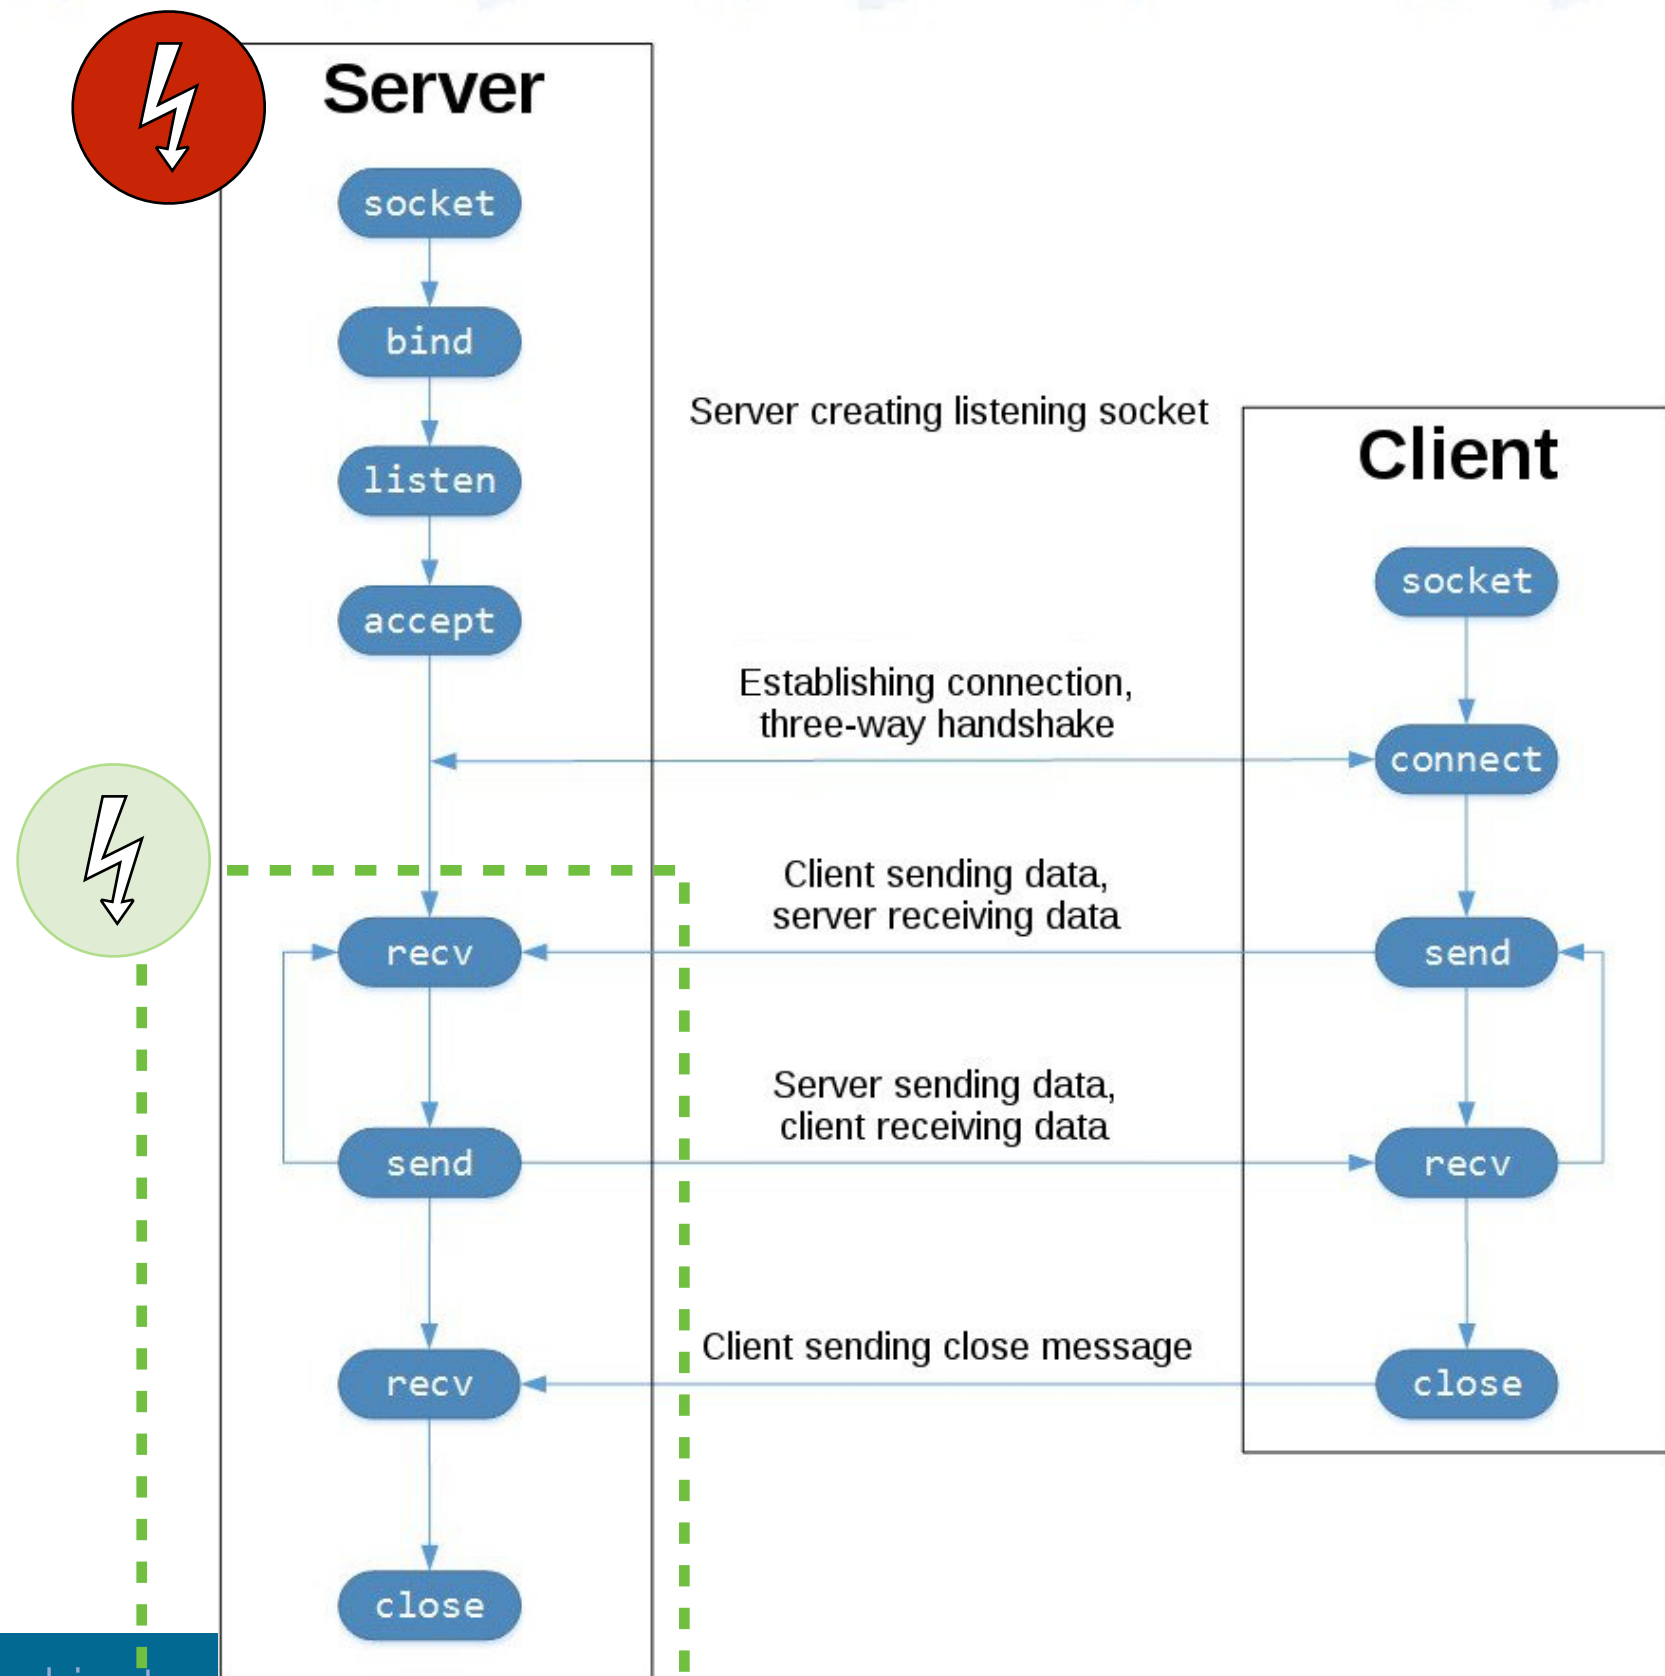
# Thread Models

How can we use threads in our Server?



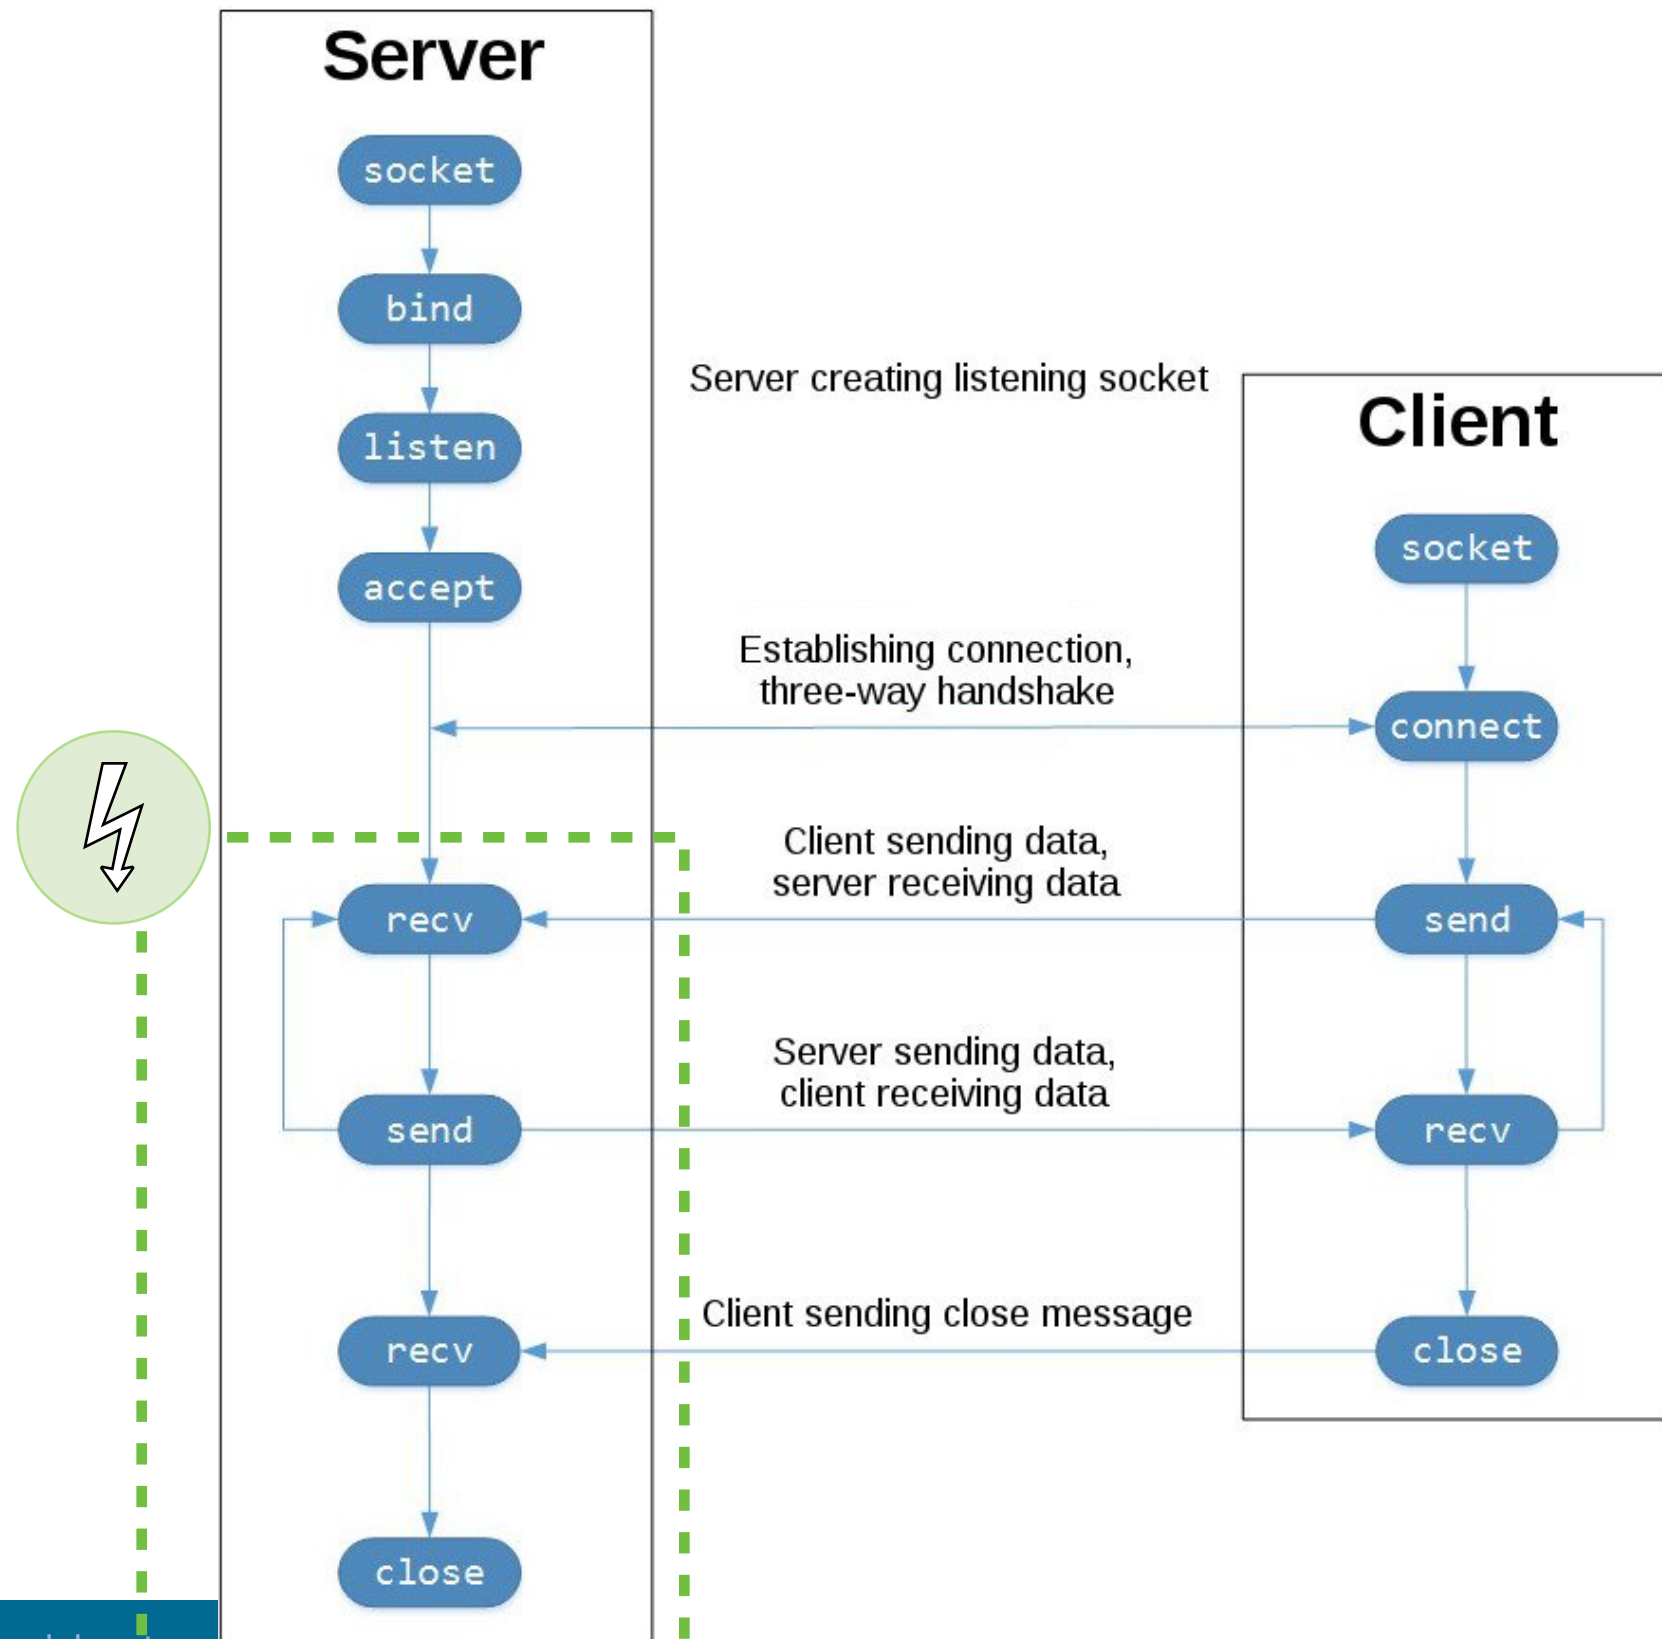
# Thread Models

How can we use threads in our Server?



# Thread Models

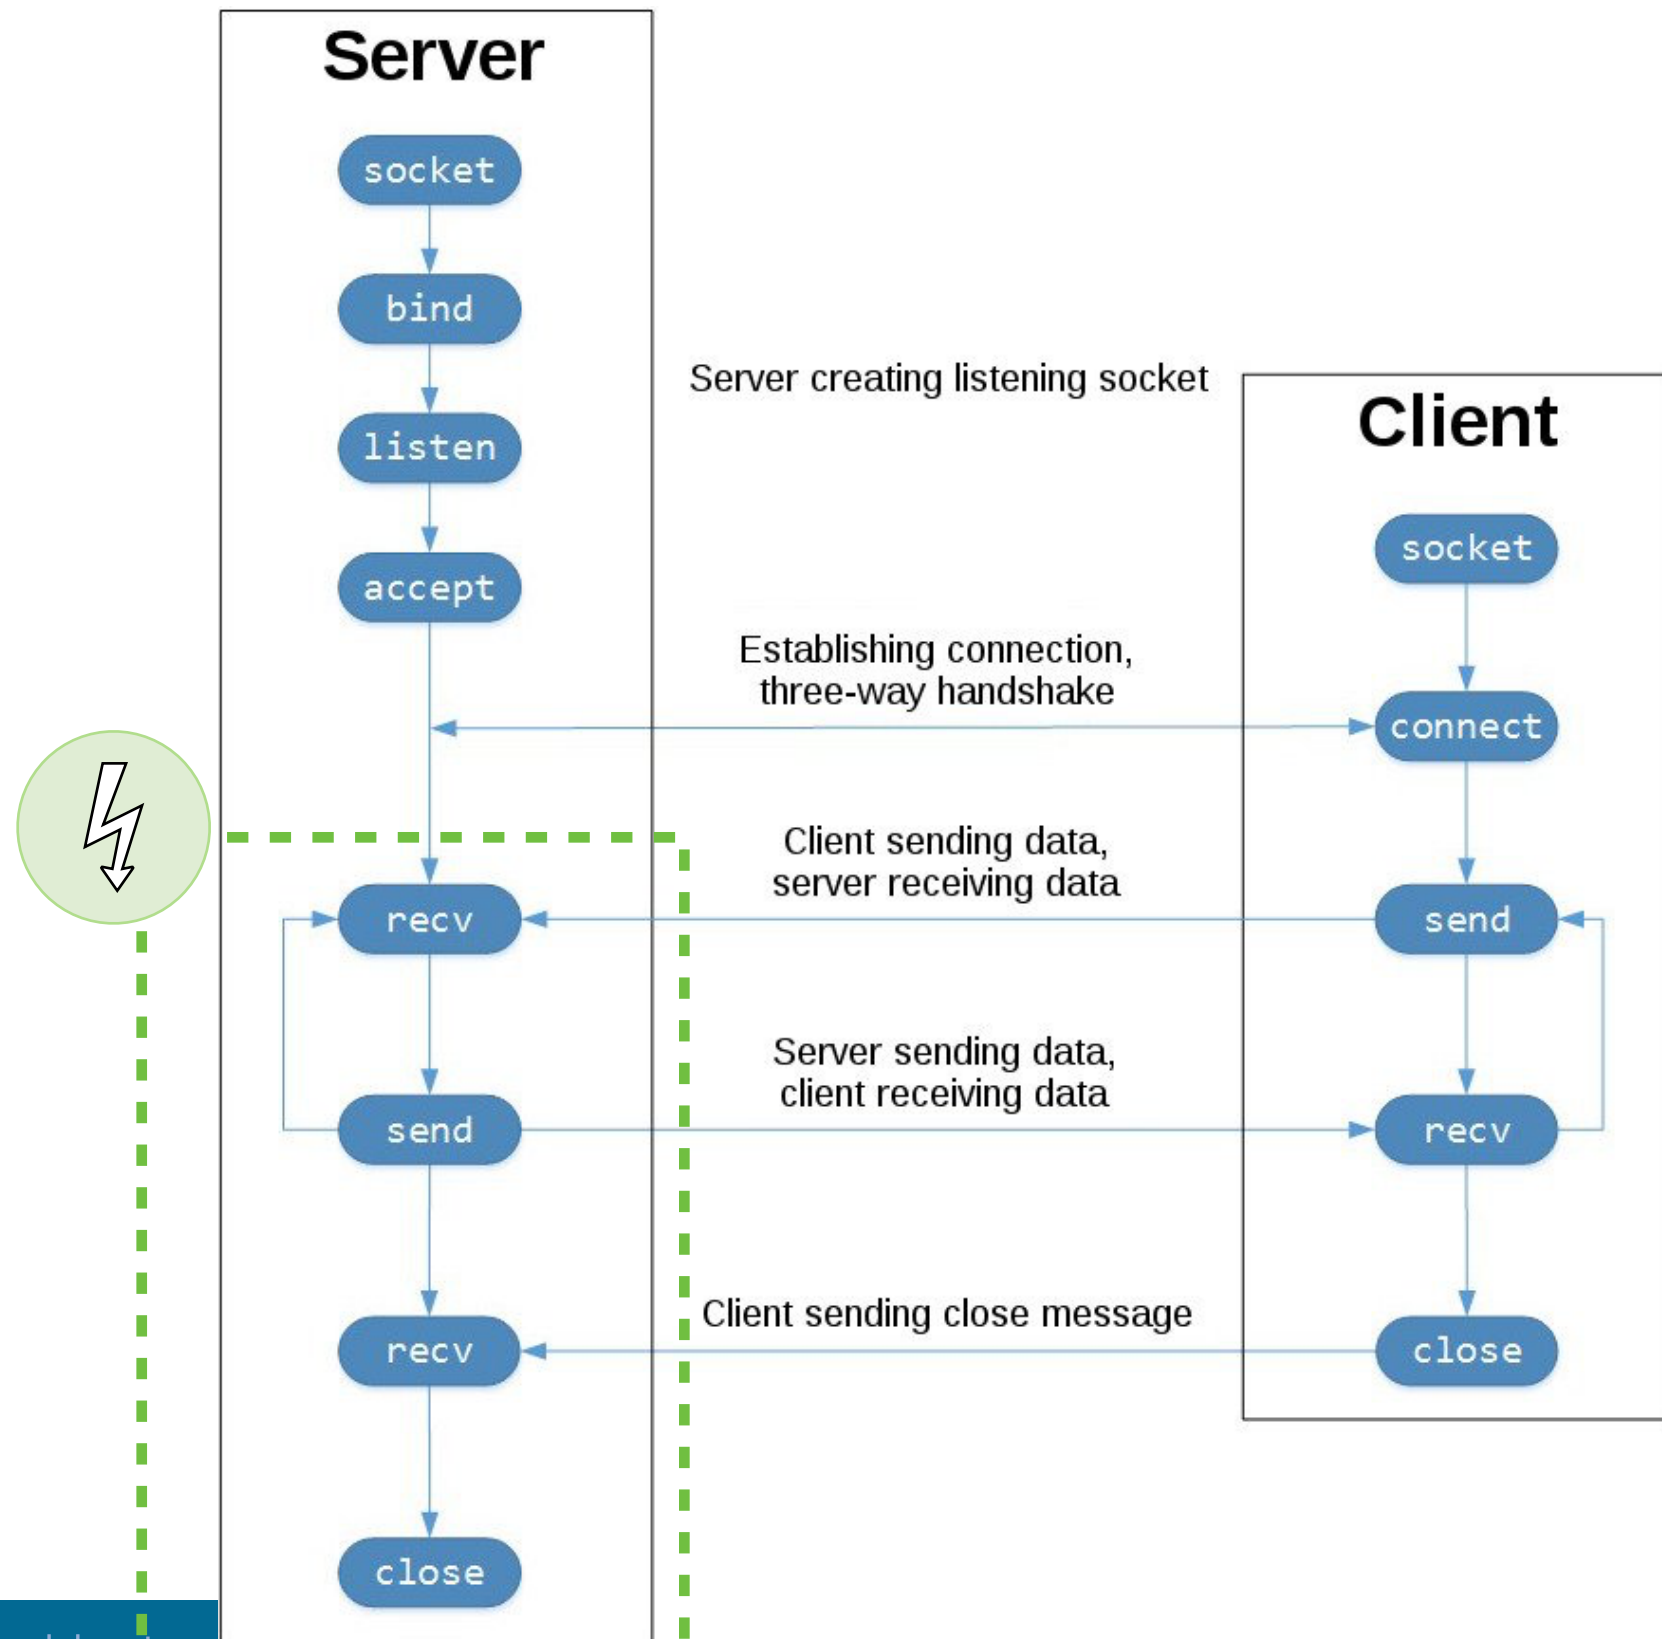
When to start threads?



# Thread Models

When to start threads?

1. On every new request create a new thread
2. When program starts create a pool of threads






# Object Pools

Common design pattern when you need to create and destroy lots of something

Create = malloc

Destroy = free

- Both of these may involve slow system calls
- Even worse if the thing you are creating is a thread!



Object pool just changes an object's state from **idle** to **in use** or back again

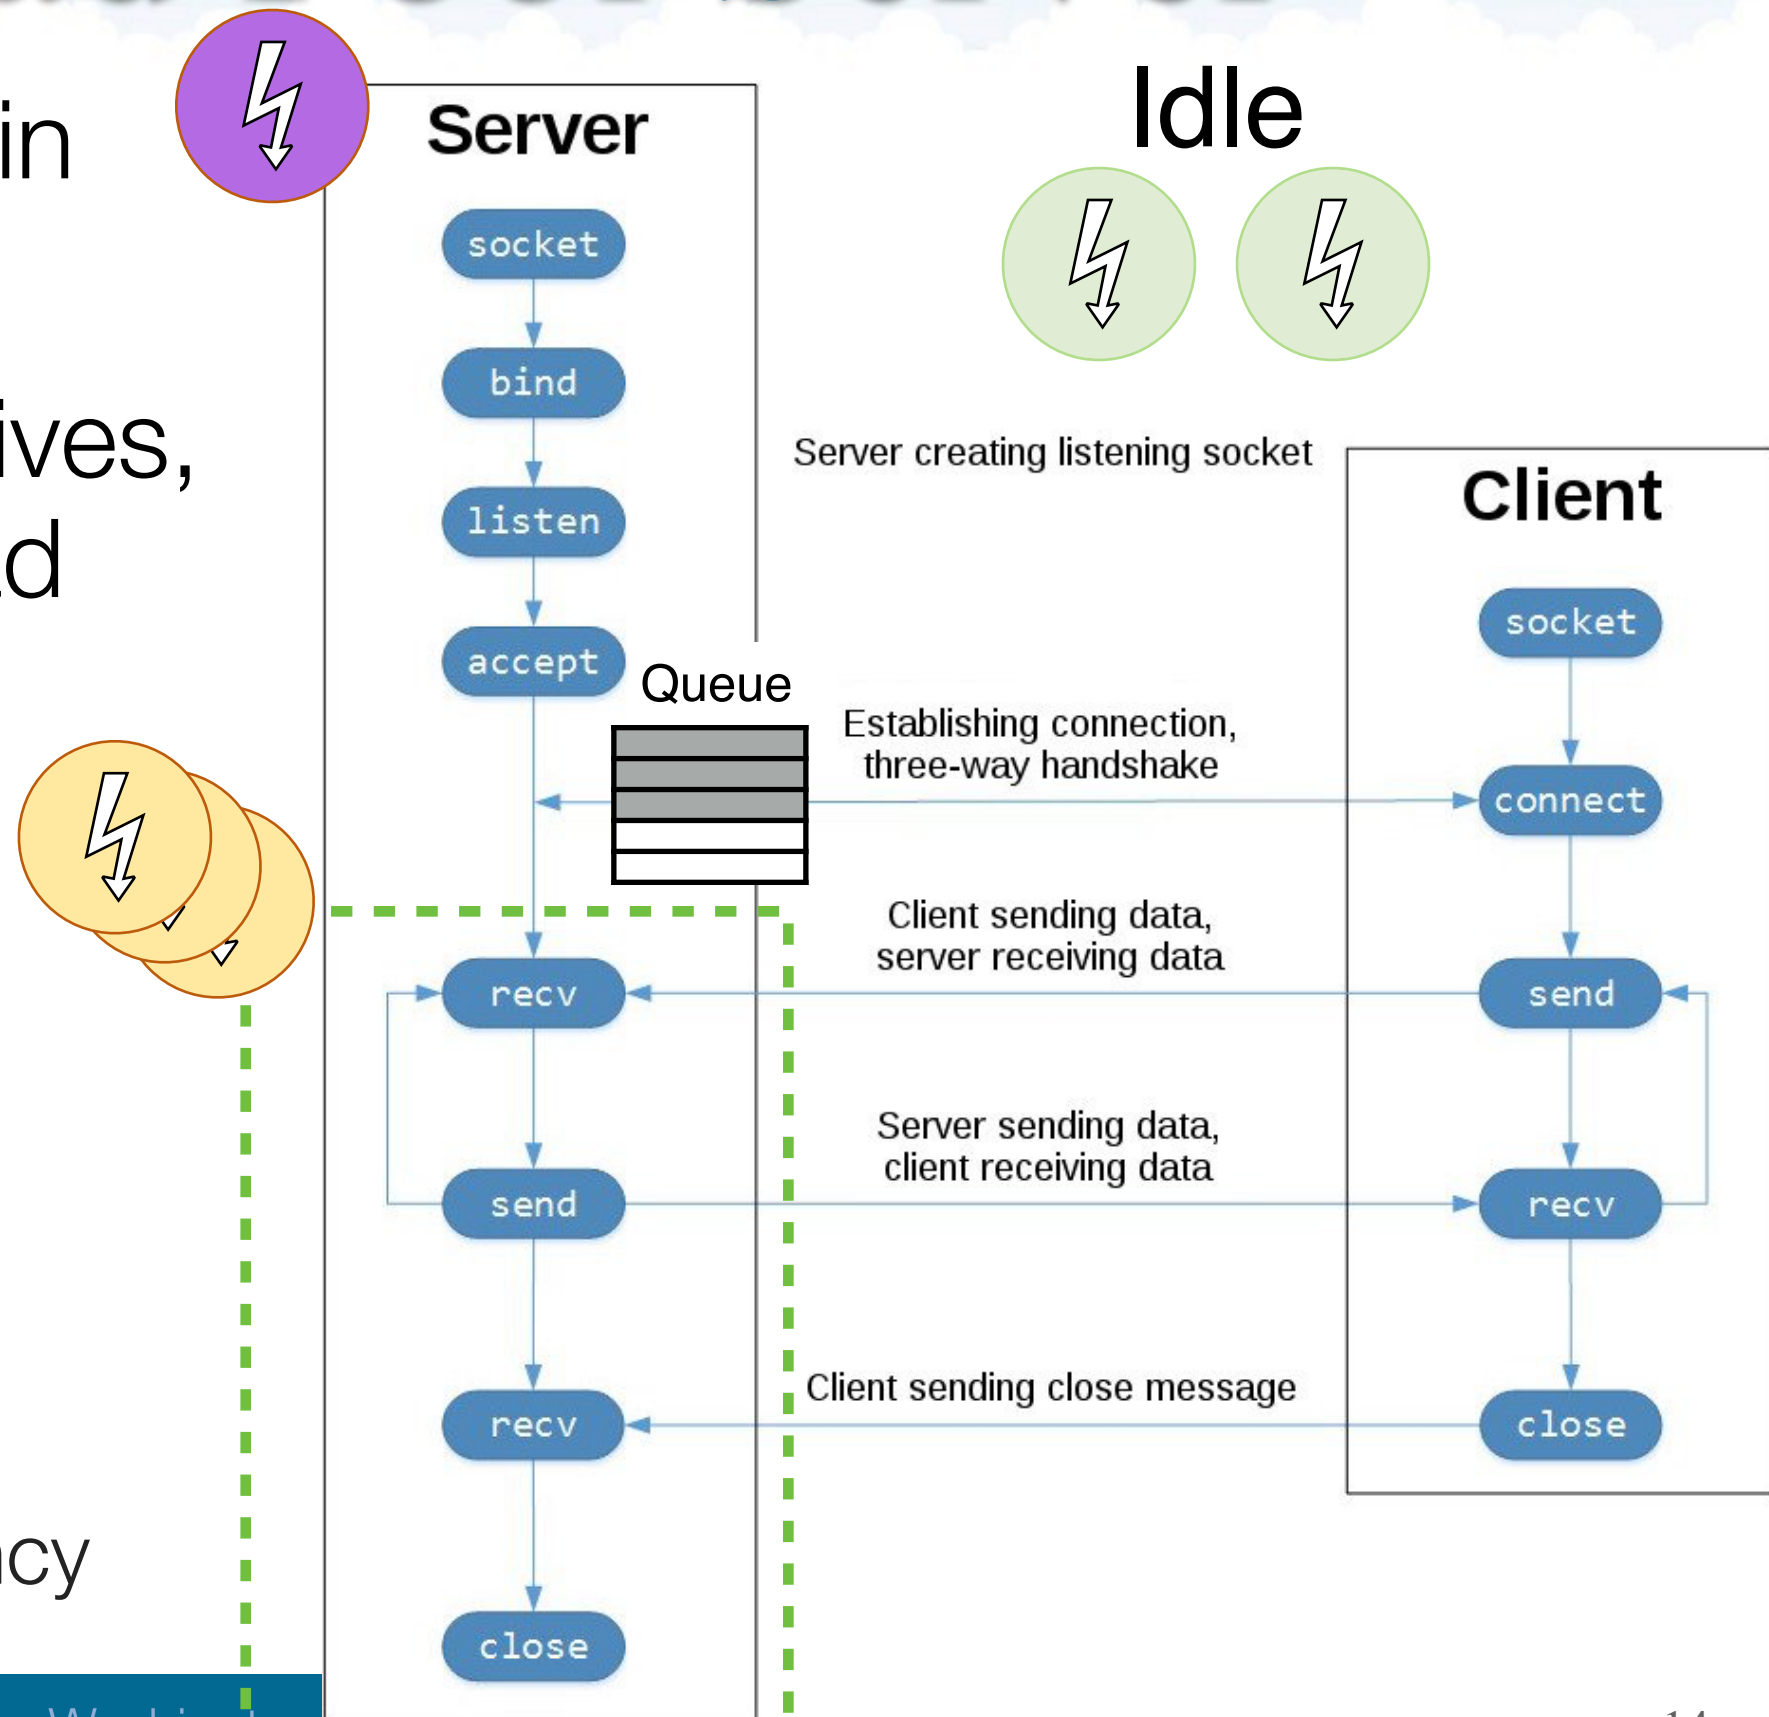
# Thread Pool Server

Idle threads wait in pool

When a client arrives, alert an idle thread

How?

- Put new client into a queue
- Wake idle thread using condition variable
- Remove client from queue using locks for consistency

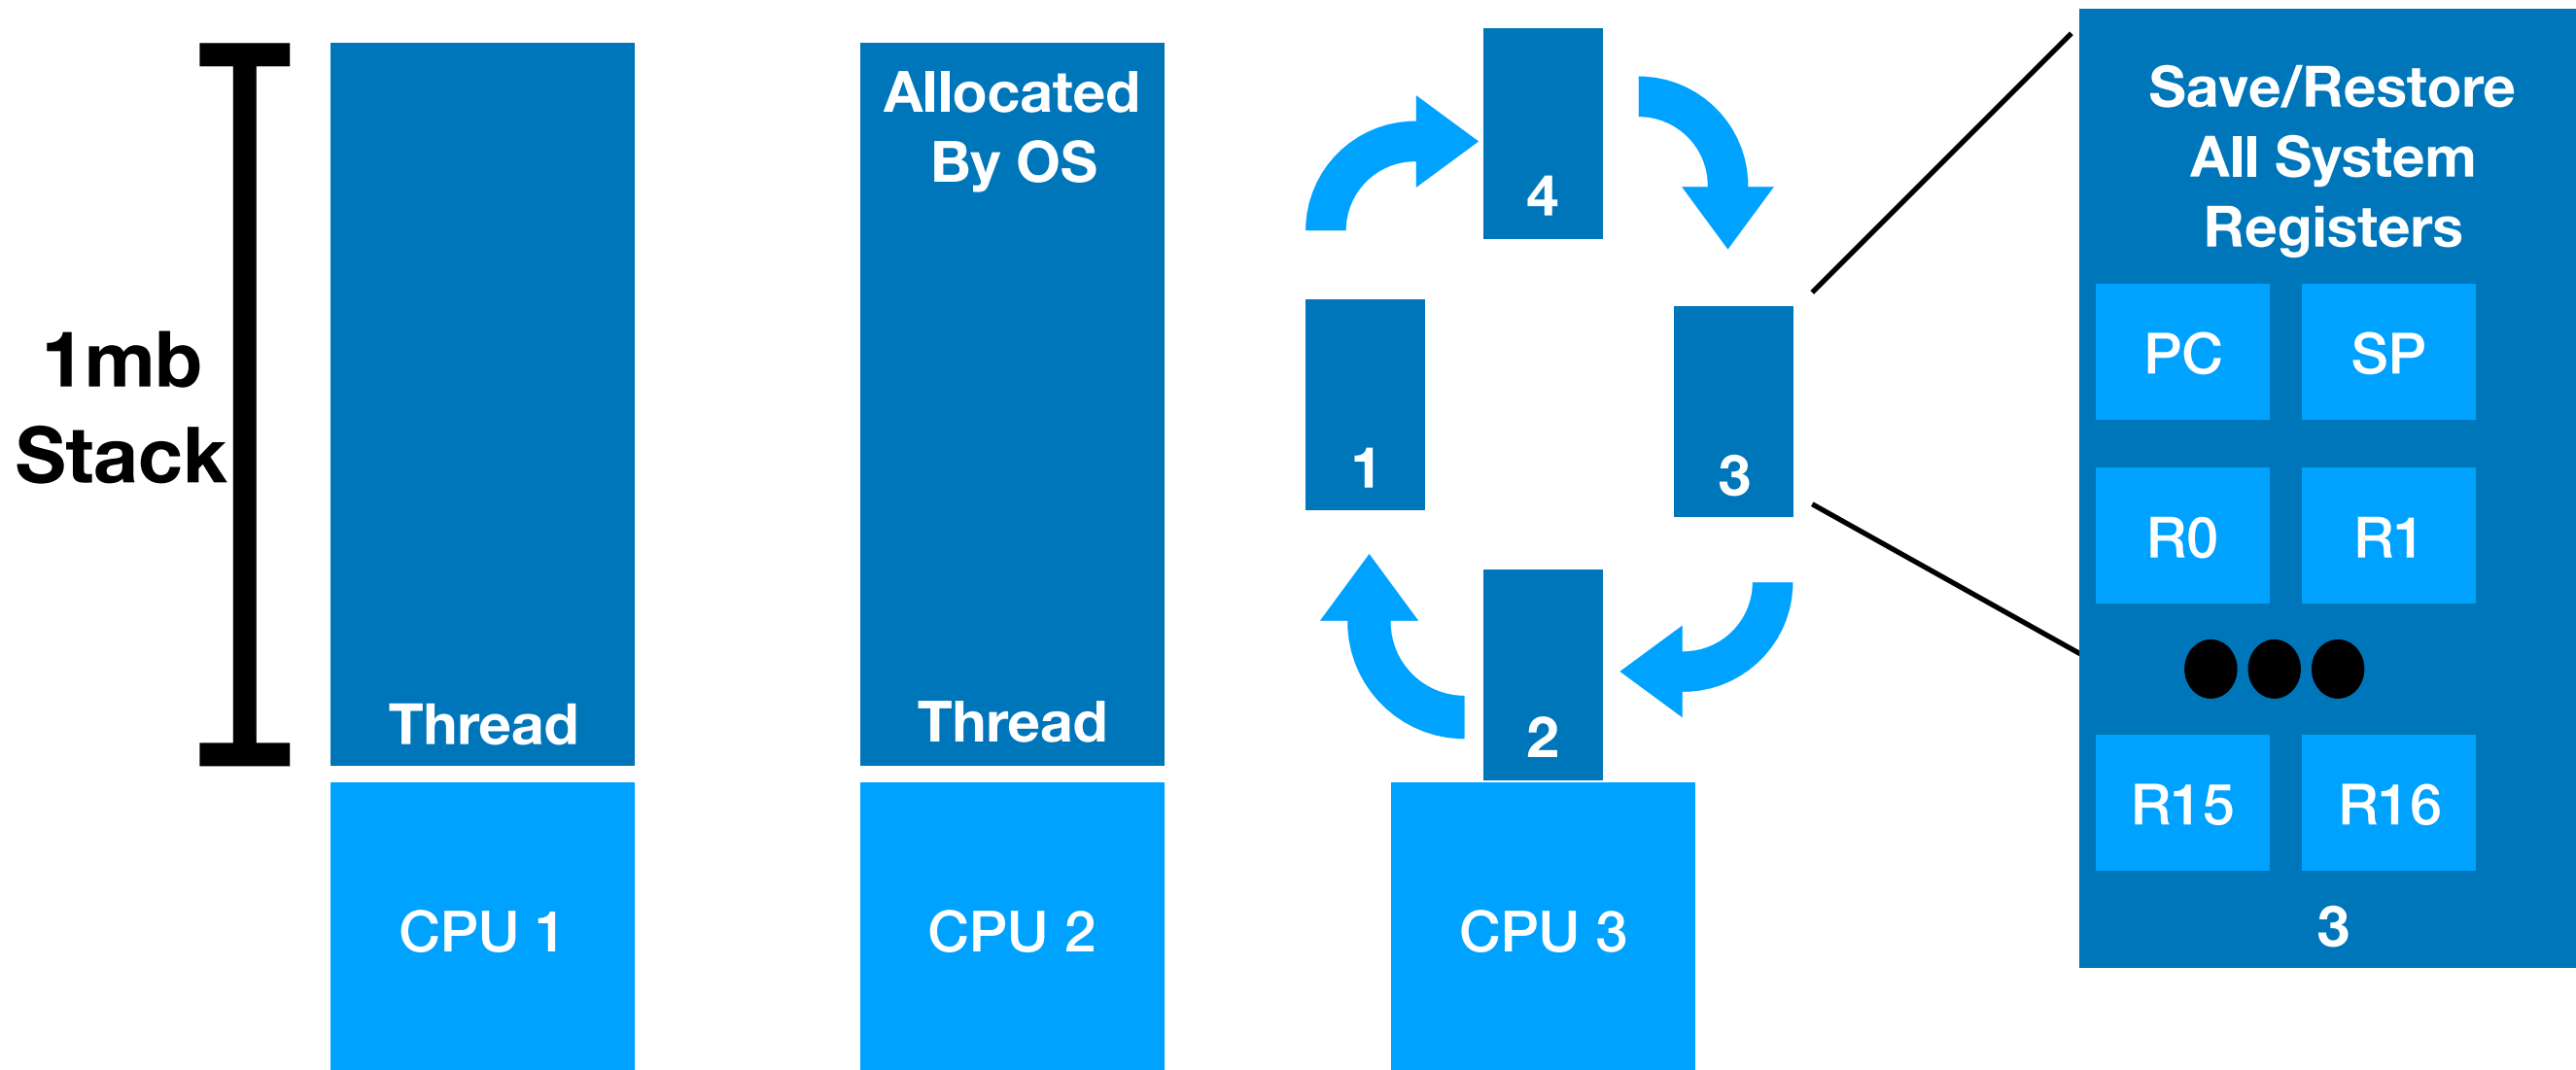




# Lightweight Threads

All about Go routines!

# A threads primer

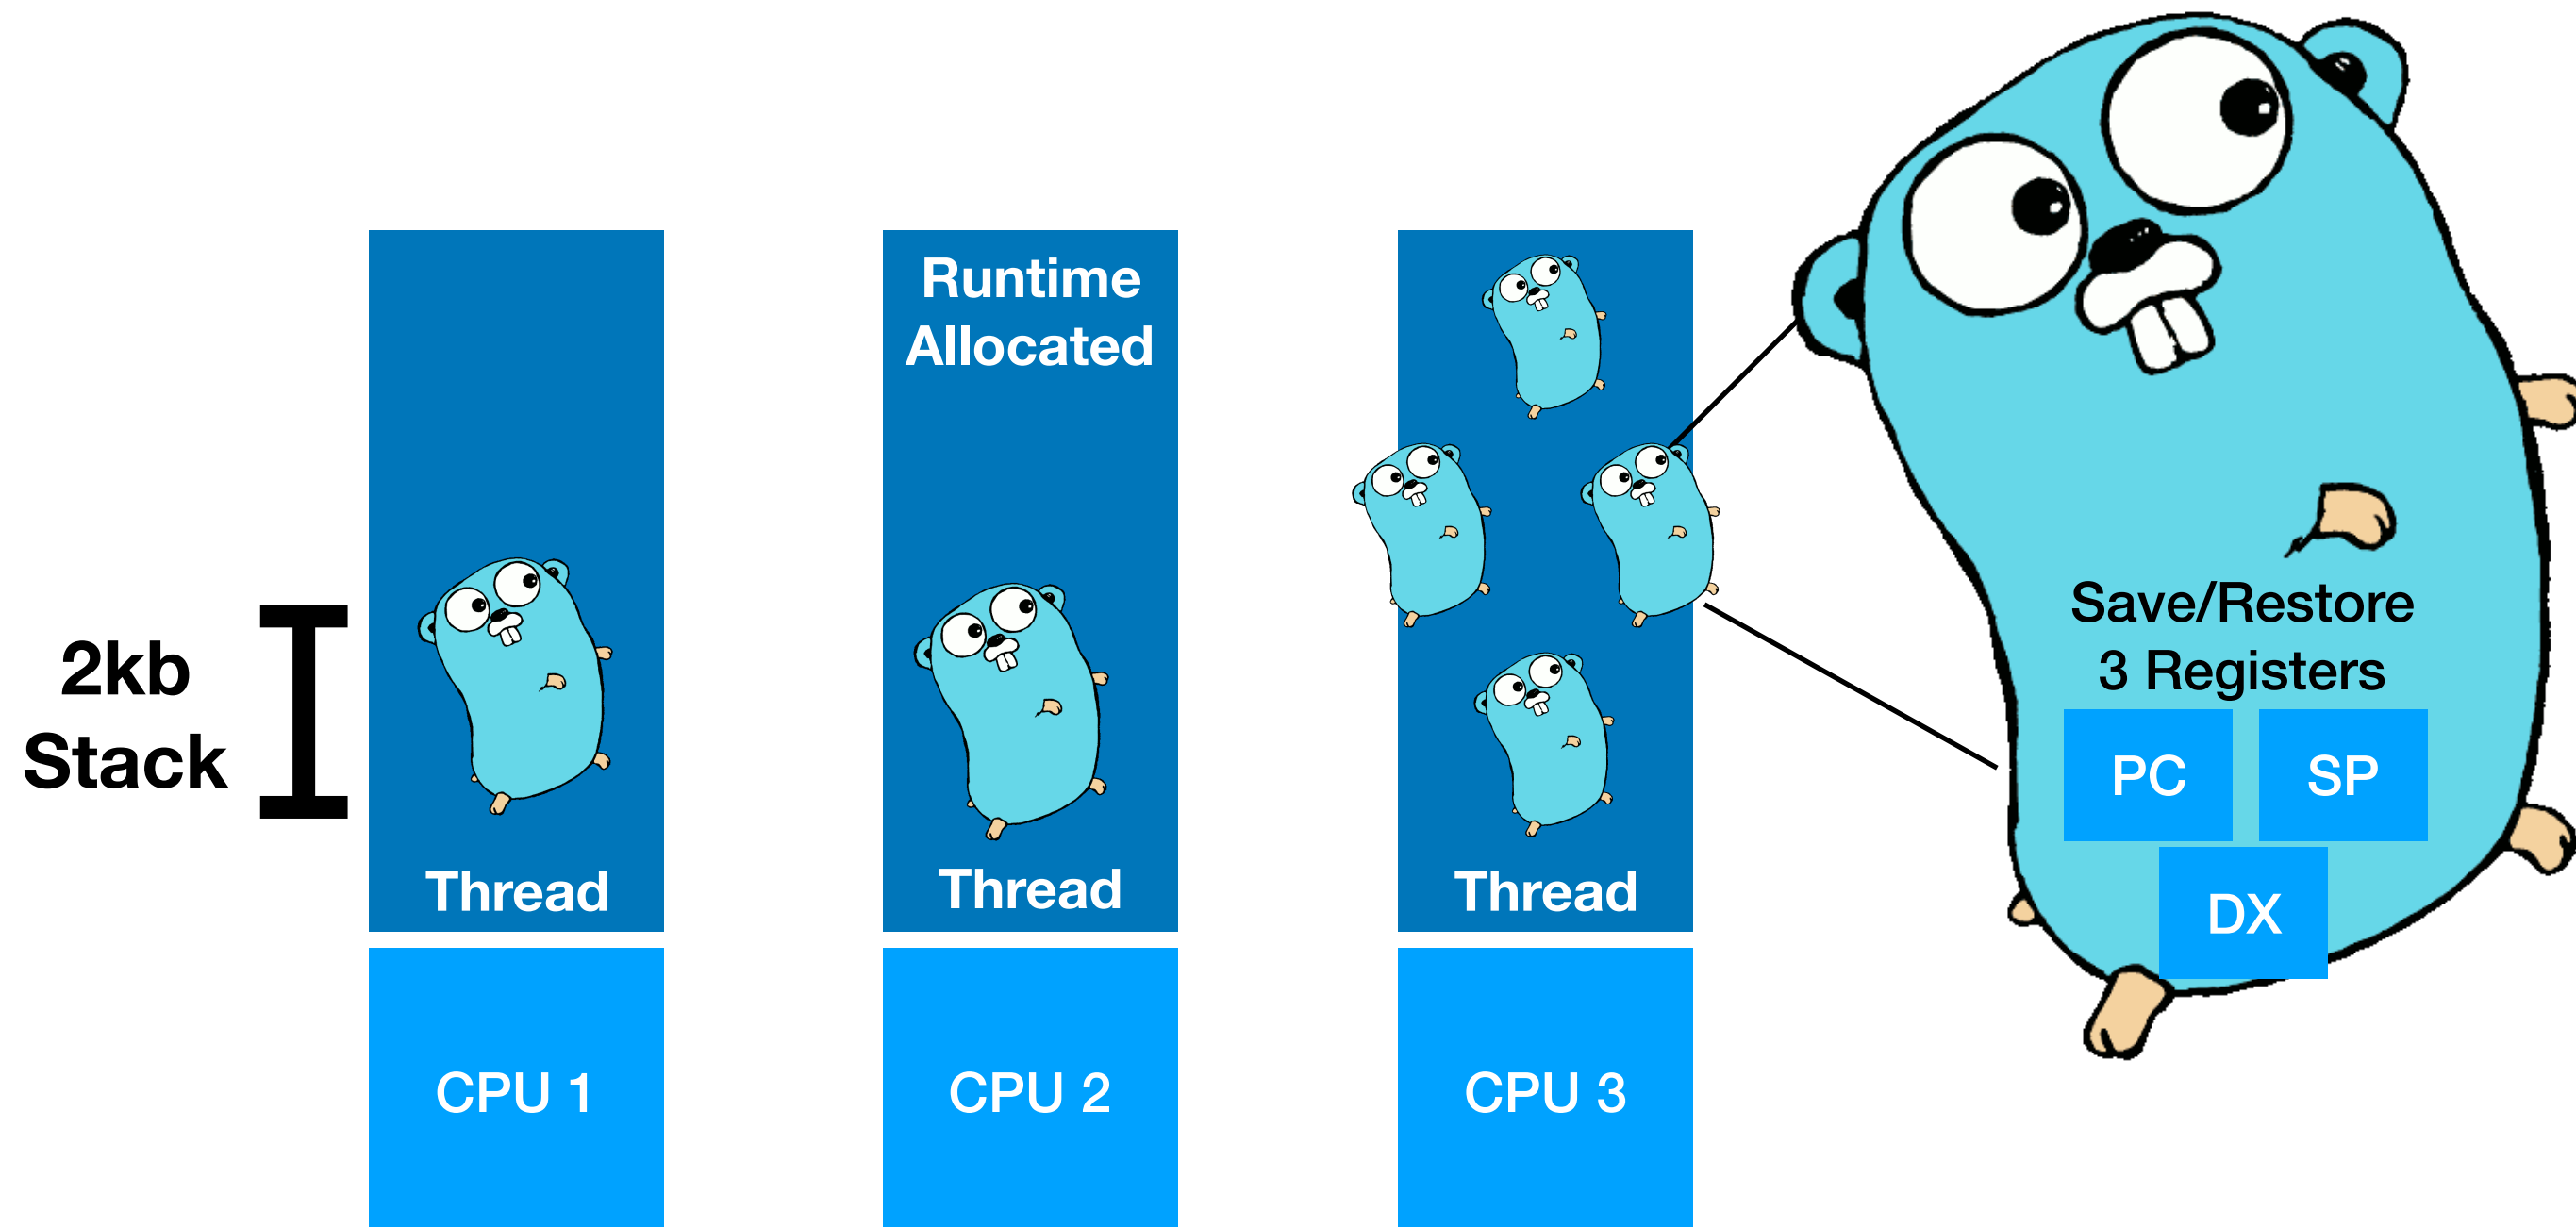


# Go Routines


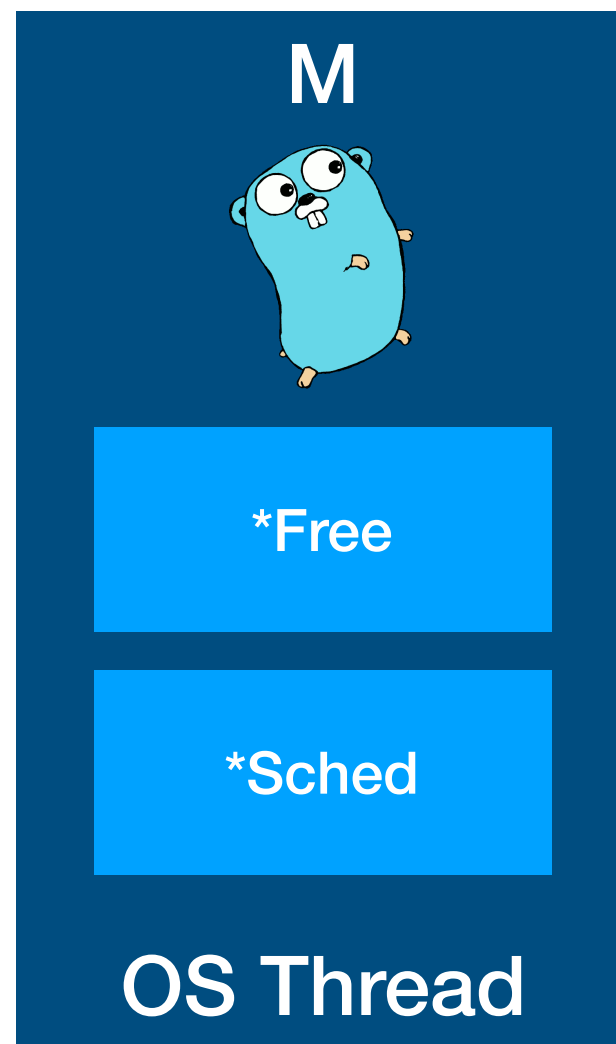
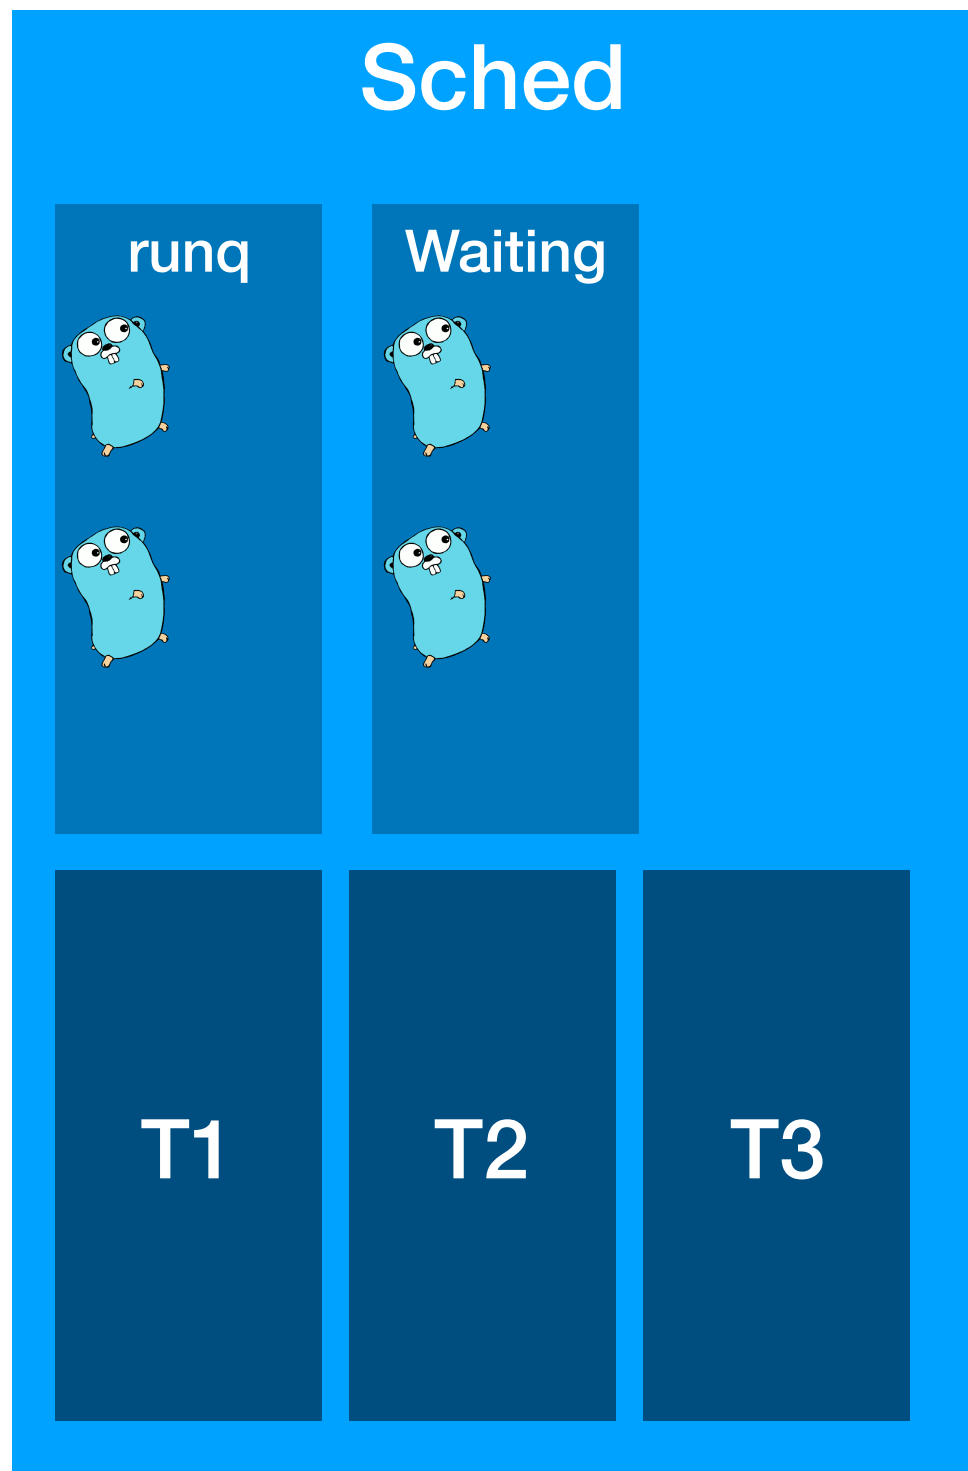
- Golang technique for concurrent programming.
- An abstraction on threading.
- Very lightweight and cheap!
- Allow programs to scale with ease

```
func helloWorld(){  
    fmt.Println("Hello World!")  
}  
  
func main(){  
    go helloWorld()  
  
    go func(txt string){  
        fmt.Println(txt)  
    }("Hello World")  
}
```

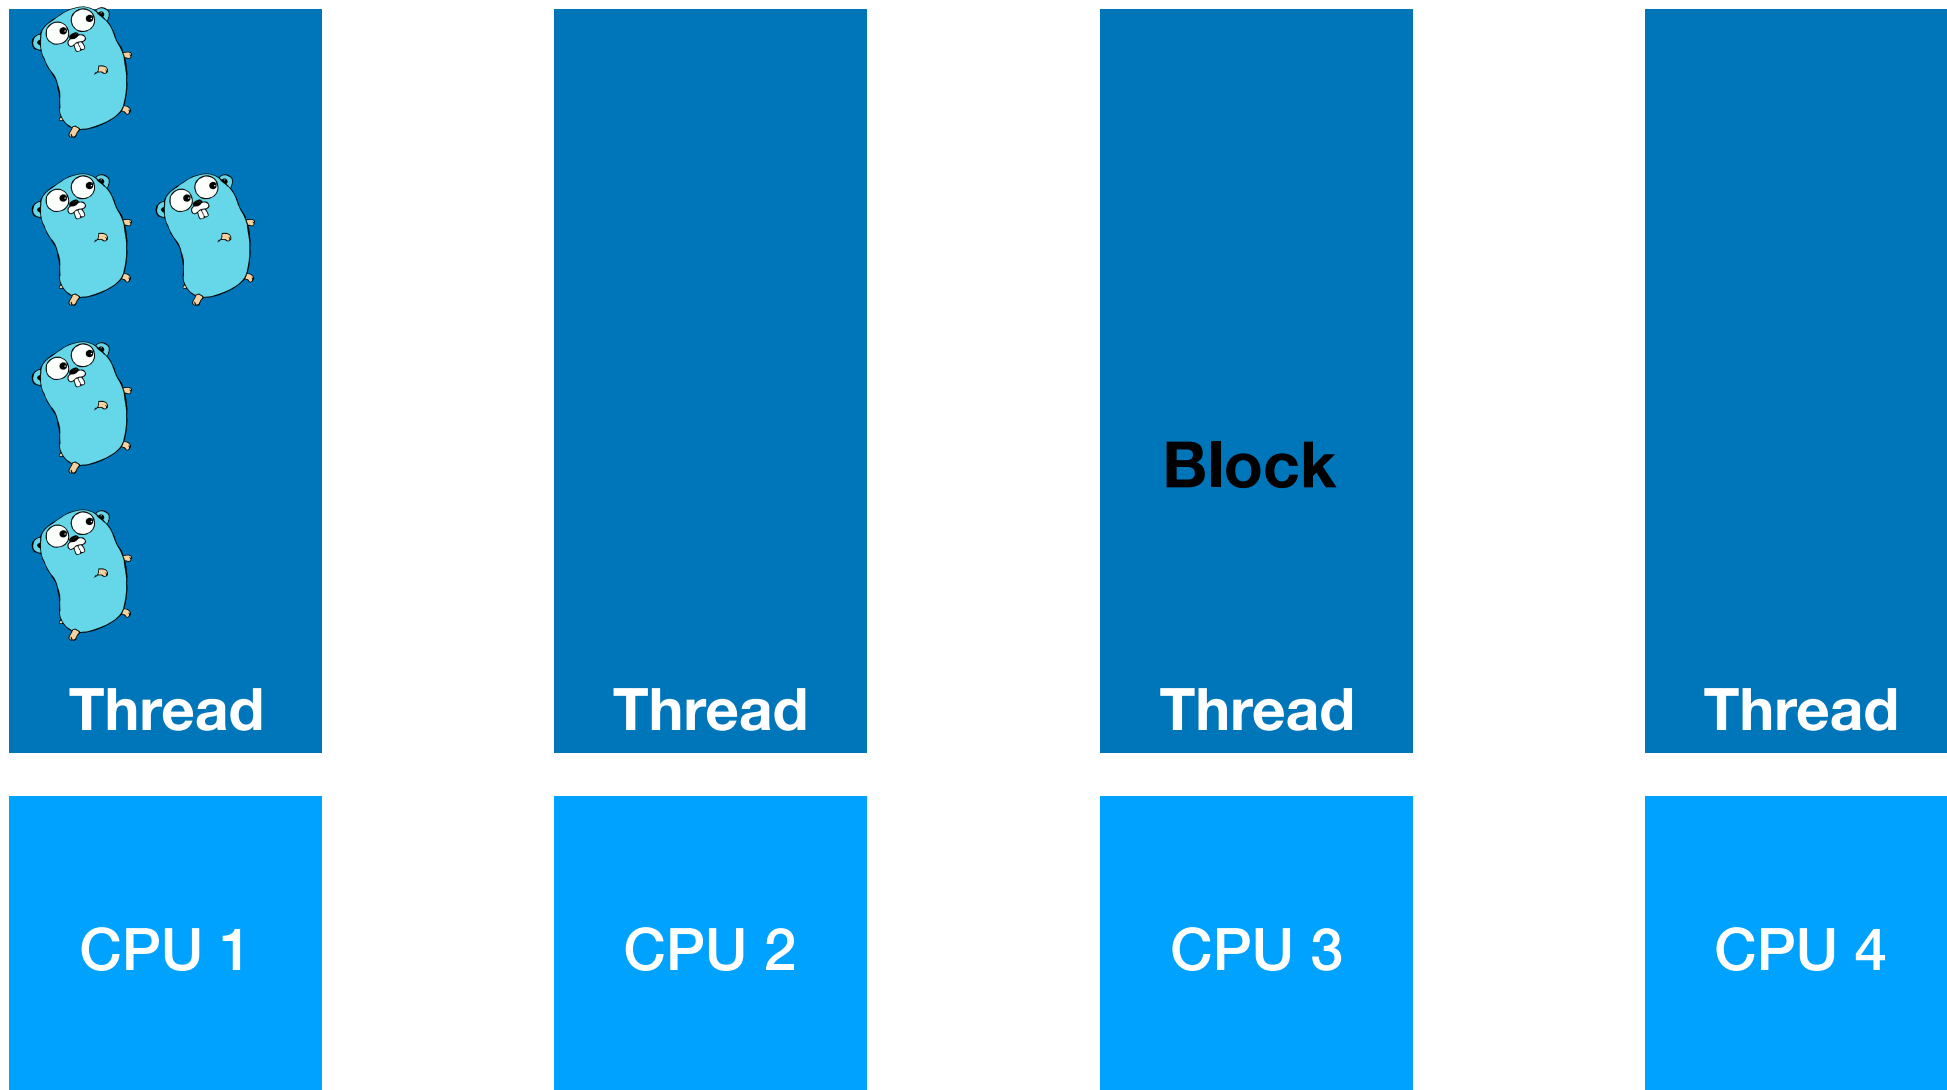
# Go Routines



# Under the Hood

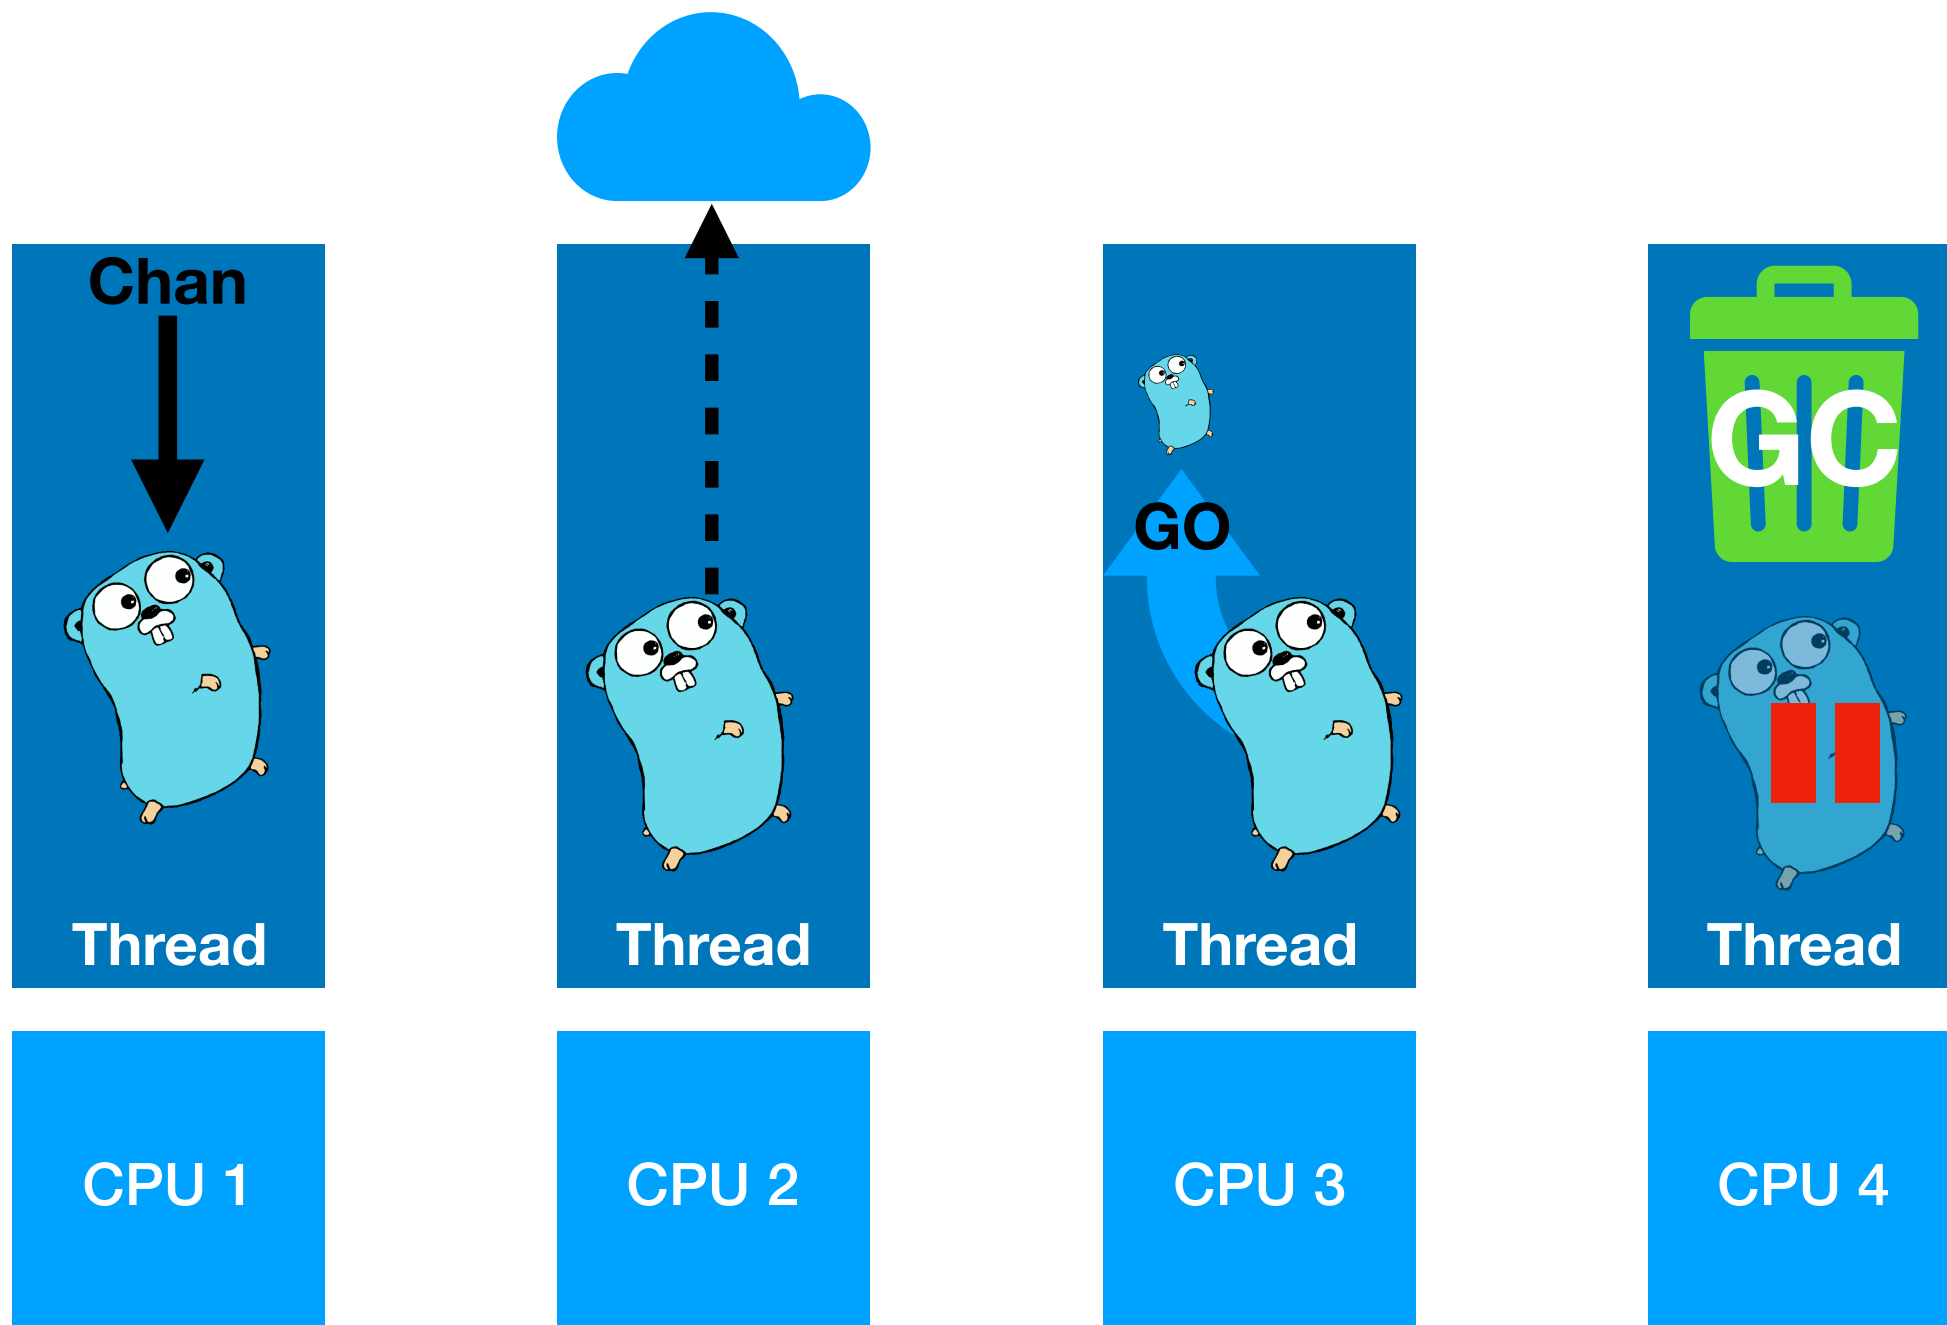


# Go Routines

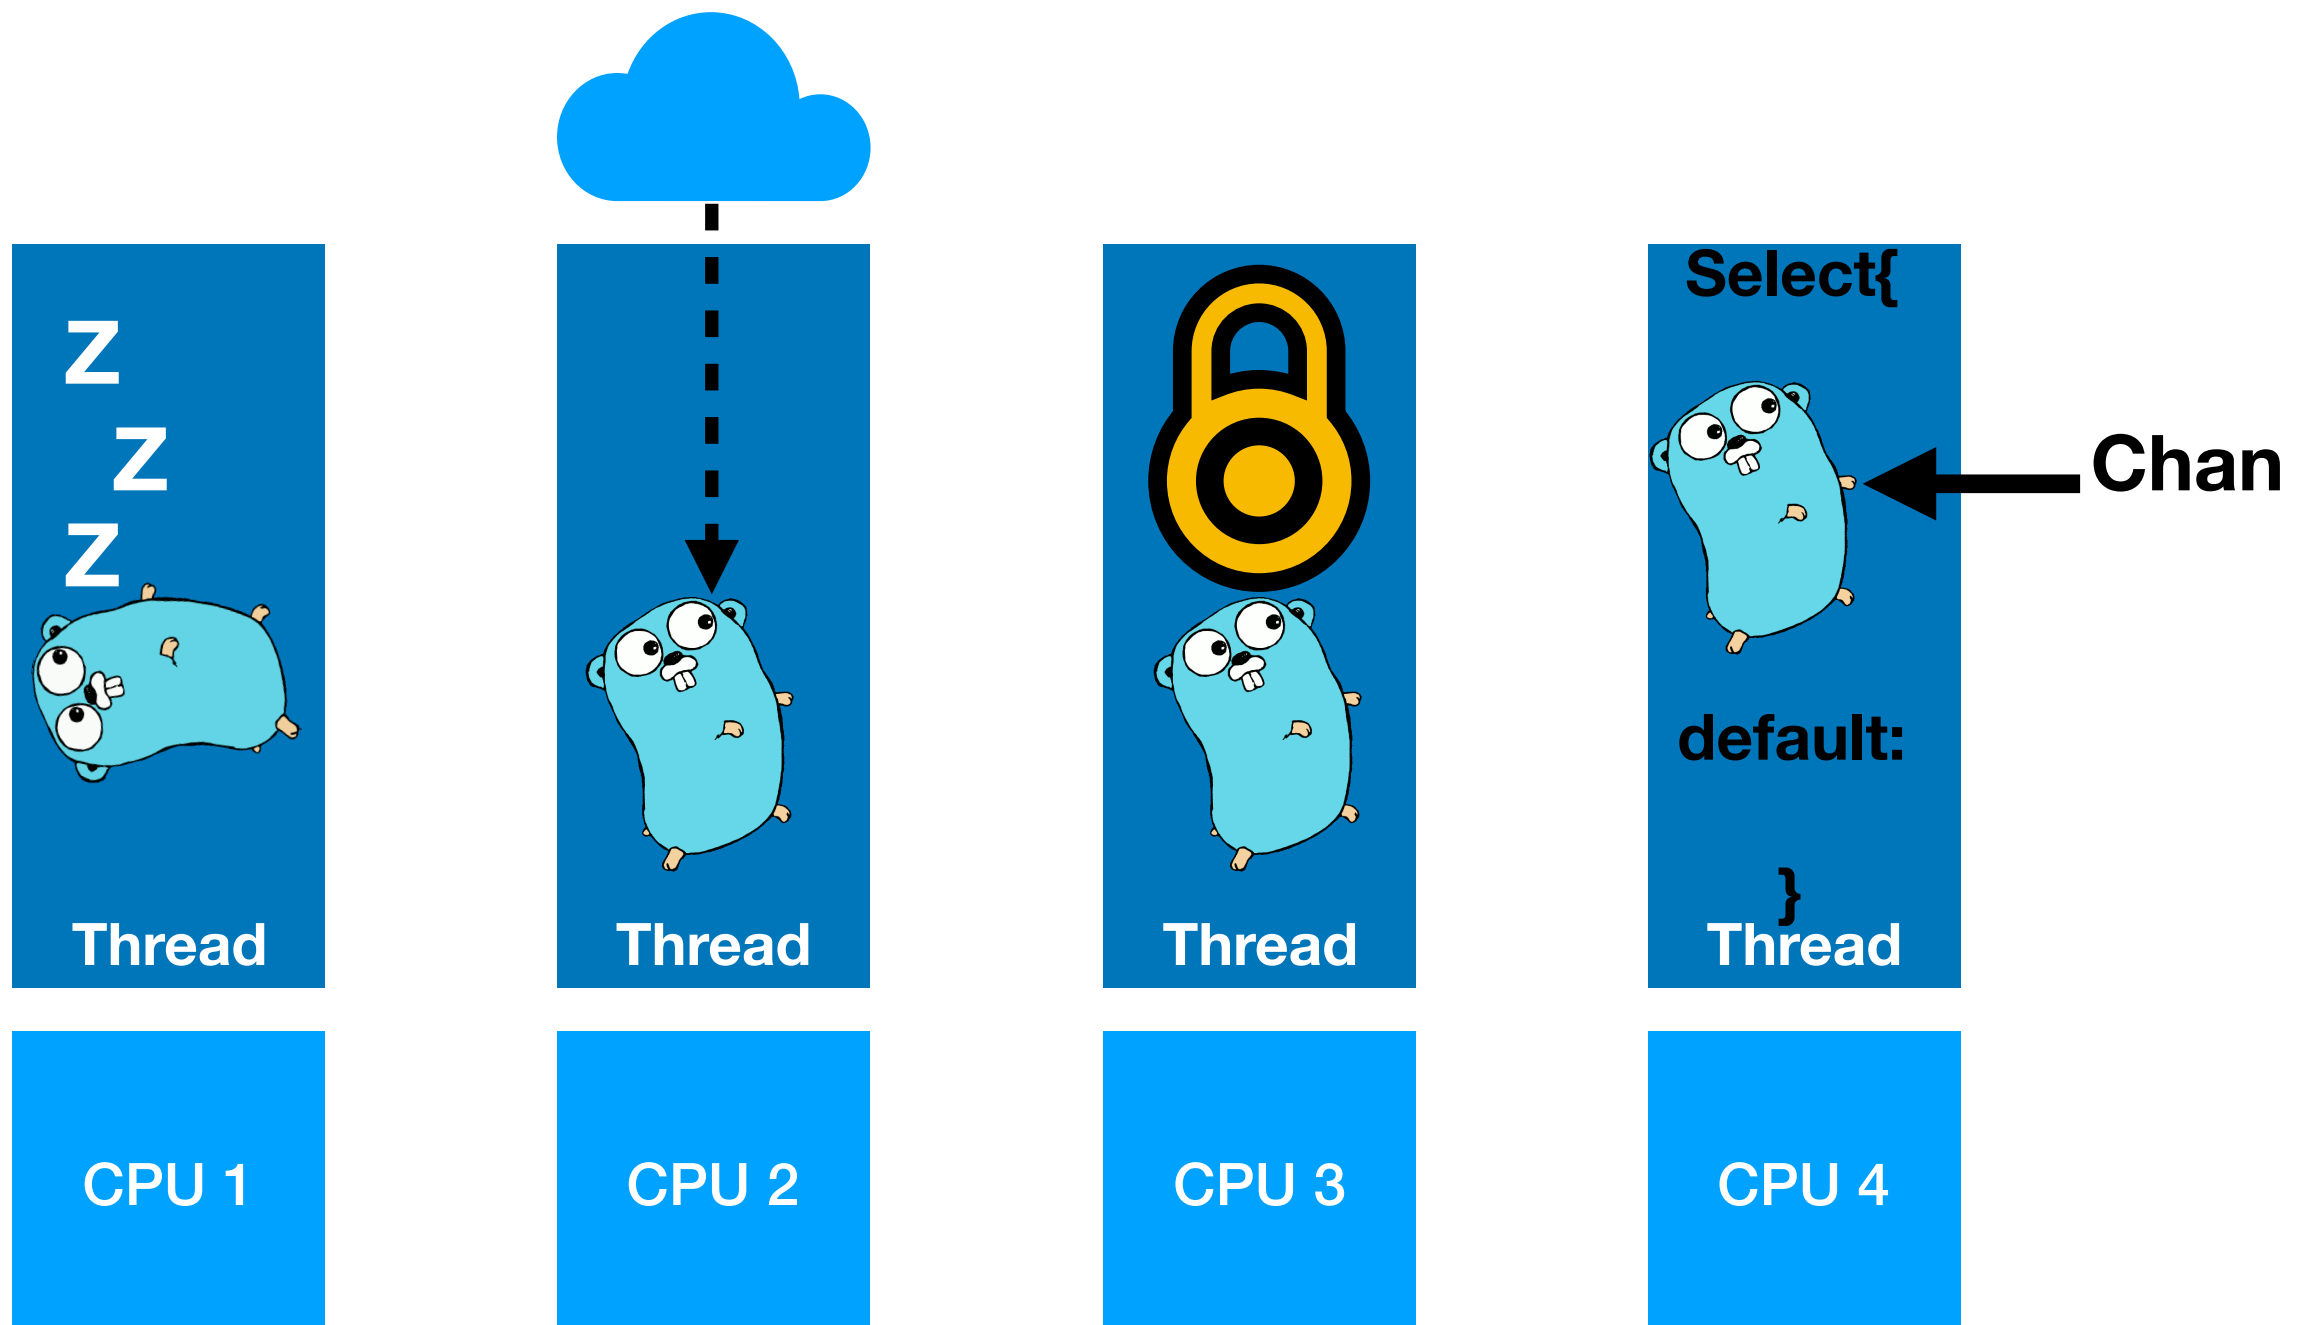




# Blocking Go Routines



# Non-Blocking Go Routines



# More threads?

Is more threads always the answer?

Threads add context switch costs and consume system resources... is there another way?

# Non-Blocking IO

Why wait?

# Blocking Calls

We needed multiple threads because recv blocks

But is it really necessary to wait on recv?

- You already saw in RUDP project that we don't need to wait forever; we can just wait for a short time and then return

Blocking / Synchronous IO:

- Go to sleep if no data, get woken up when it arrives

Non-Blocking / Asynchronous IO:

- Check if there is data, do something else if no data, check again

# Simple Non-Blocking

Sockets can be set to non-blocking mode

```
import socket
# Create a TCP/IP socket in non-blocking mode
server = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
server.setblocking(0)
```

Then **recv** calls will not wait for data, just return error

```
while True:
    try:
        data = conn.recv(1024)
    except socket.error:
        print("No data yet")
```

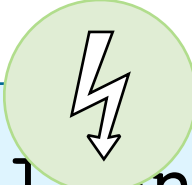
Drawbacks of this approach?



# Non-Blocking Server

What happens if we have many clients?

Client 1   
Client 2   
Client 3   
...  
Client n 



```
# Accept all clients...  
  
for client in clients:  
    try:  
        data = client.recv(1024)  
        process(data)  
    except socket.error:  
        print("No data yet")
```

Code is messy and inefficient if many clients!

# Non-Blocking IO

We need a better way to know what data is ready!

## **select** event polling

- Register a set of IO “file descriptors” you care about
- Sleeps until at least one of them has data -> won't block!

```
int select(int nfd, fd_set *readfds, fd_set *writefds,  
          fd_set *errorfds, struct timeval *timeout);
```

- Assumes a Unix environment where files, sockets, and other types of IO are all mapped to a file interface

# Select Example

```
import selectors
import socket
```

```
def accept(sock, mask):
    conn, addr = sock.accept()
    conn.setblocking(False)
    sel.register(conn,
                  selectors.EVENT_READ, read)
```

```
def read(conn, mask):
    data = conn.recv(1000)
    if data:
        conn.send(data)
    else:
        sel.unregister(conn)
        conn.close()
```

```
sel = selectors.DefaultSelector()
sock = socket.socket()
sock.bind(('localhost', 1234))
sock.listen(100)
sock.setblocking(False)
sel.register(sock, selectors.EVENT_READ, accept)

while True:
    events = sel.select()
    for key, mask in events:
        callback = key.data
        callback(key.fileobj, mask)
```

# Non-blocking Variants

Languages, runtimes, and OS's typically have several ways to do non-blocking IO

**select:** system call for checking if things are ready

**epoll / kqueue:** app/OS interface for checking if things are ready (much more efficient than original select)

But now select can be viewed as an API, and might be implemented with something like epoll.

# Event-Based Programming

Registering call backs for events can be a simpler programming model

- Simpler to write... maybe harder to debug!

Adds a layer of abstraction

- Event notification layer checks for events and decides what order to process them in. Why is this helpful/interesting?
- Could use multiple threads to process the events!

# node.js

Web framework for javascript-based apps

Probably the most popular event based platform

Single threaded event based server!

- Faster and less resource intensive than many multi-threaded servers!

Other event based frameworks/languages:

- Erlang, Elixir, ...



# Assignment 2

# Technical Writing

Being able to present ideas is just as important as being able to write code!

[ ] Write a blog post on a networking topic

- Must be long enough to be interesting
- You must write some code or run experiments
- Present useful information in an understandable way
- Present useful information in a visually appealing way

# Ideas

## Performance comparison of...

- Node.js vs Apache vs nginx vs ...
- Thread pool vs new thread per request in language X
- http vs https vs http2

## Tutorial on...

- how to use wireshark to analyze HTTP traces or solve a puzzle
- how to gather statistics of public wifi traffic (ethically)
- how to use go co-routines and how they work under the hood
- queueing theory 101 with example measurements
- how to use epoll / select / etc in language X
- everything that happens when you open a page in a browser
- python 2 vs python 3 networking code
- how to generate traffic to benchmark a web server

# Inspiration

Julia Evans' blog and zines

- <https://jvns.ca/>